

# CS-300: Data-Intensive Systems

## The Storage Layer

(Chapter 12.1-12.5    13.5)


*Prof. Anastasia Ailamaki, Prof. Sanidhya Kashyap*



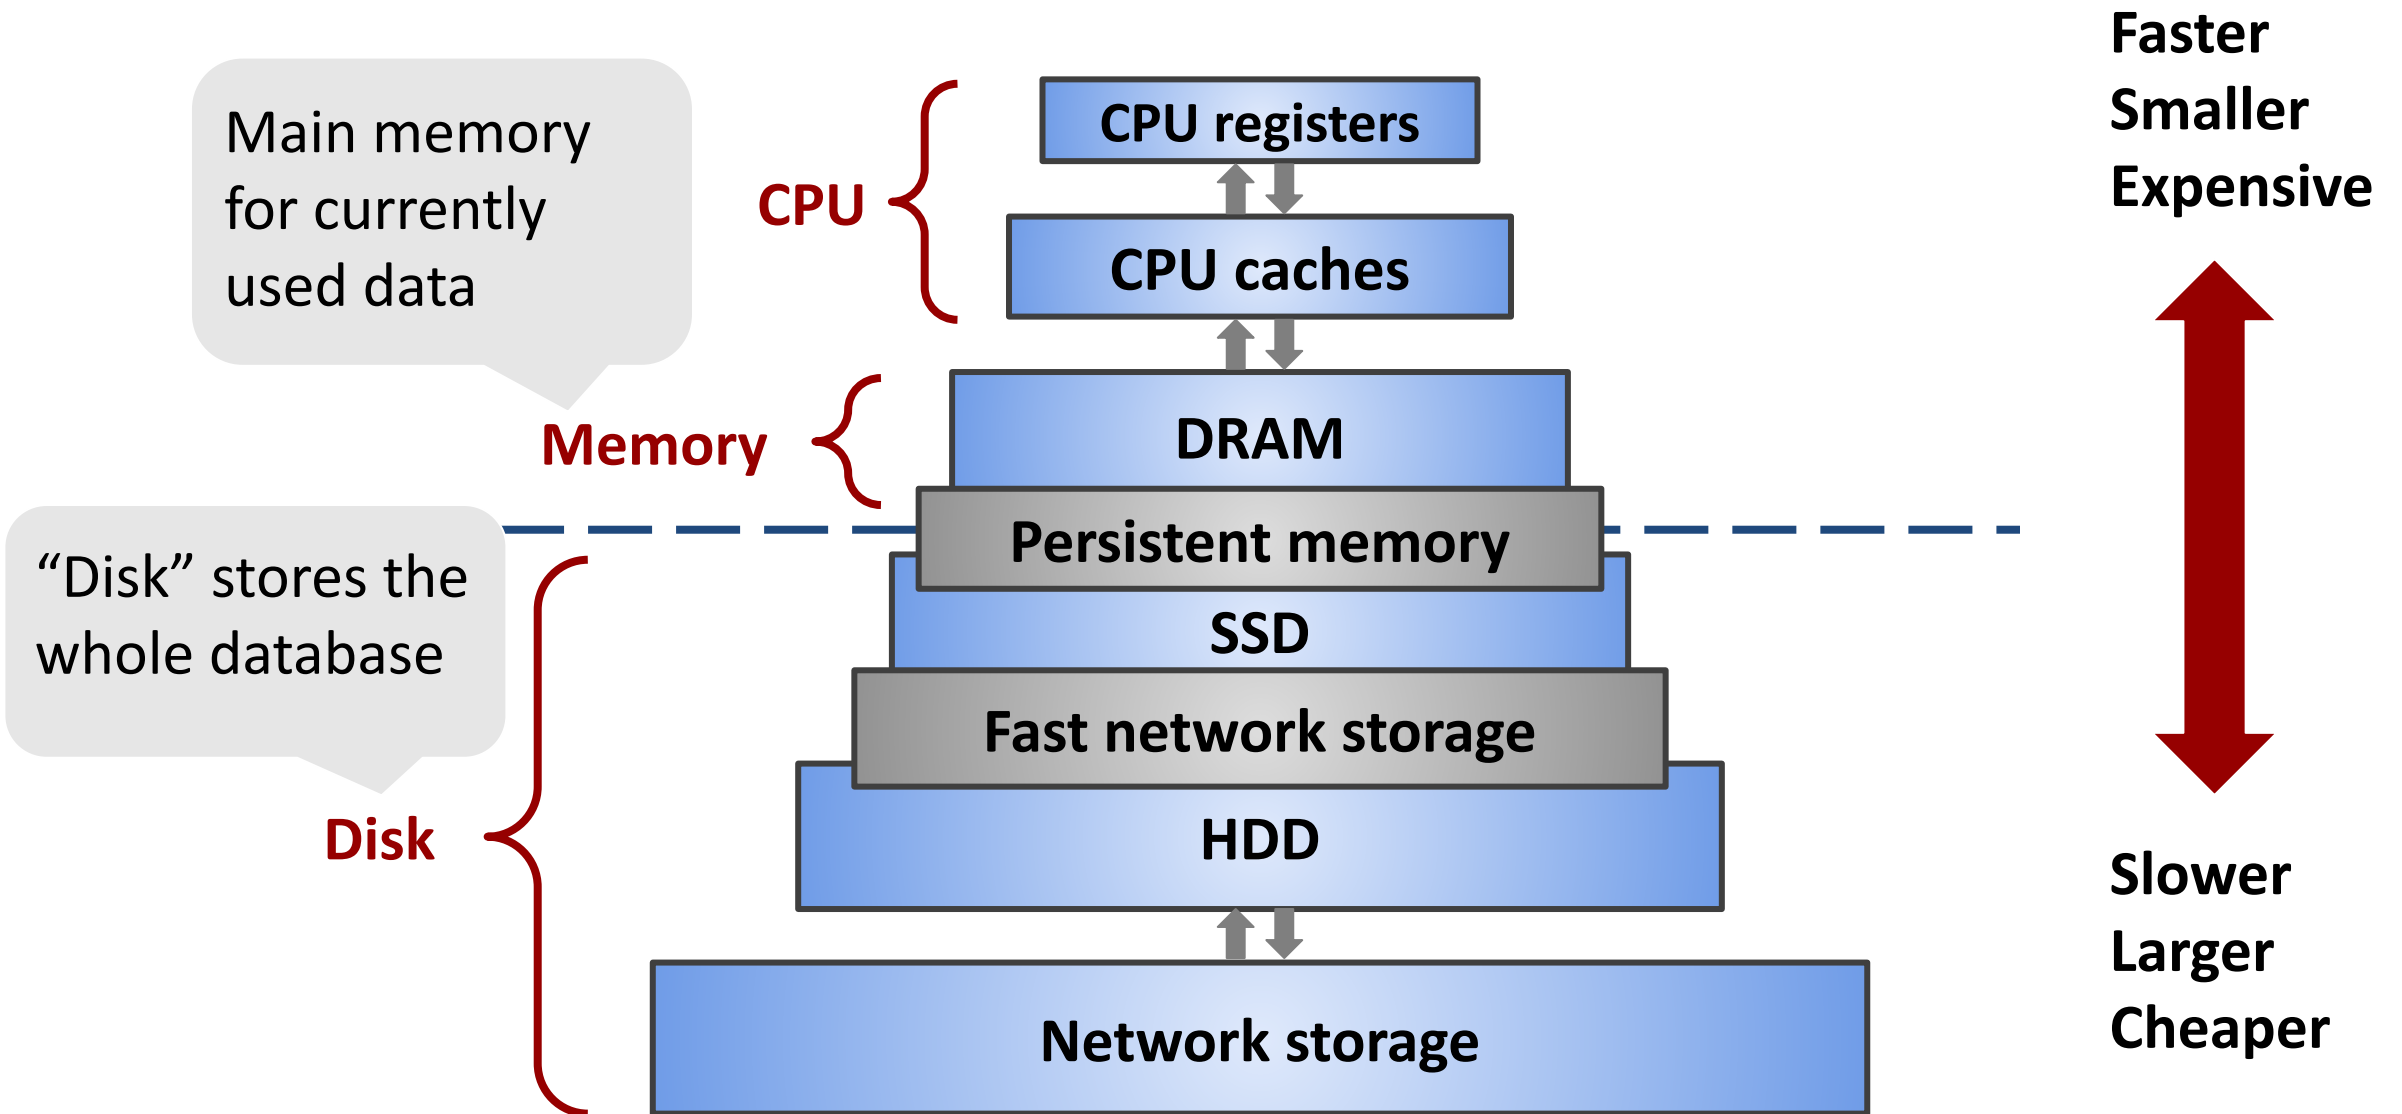
# Today's focus

- DBMS layers and storage hierarchy
- Disks
  - HDD
  - SSD
- Buffer management








# The storage hierarchy



# The storage hierarchy



# Latency numbers every programmer should know

<b>1 ns</b>	L1 cache reference	 <b>1 sec</b>
<b>4 ns</b>	L2 cache reference	 <b>4 sec</b>
<b>100 ns</b>	DRAM	 <b>100 sec</b>
<b>16,000 ns</b>	SSD	 <b>4.4 hours</b>
<b>2,000,000 ns</b>	HDD	 <b>3.3. weeks</b>
<b>~50,000,000 ns</b>	Network storage	 <b>1.5 years</b>
<b>1,000,000,000 ns</b>	Tape archives	 <b>31.7 years</b>

# Disks and Files

- DBMS stores information on disks
  - In an electronic world, disks are a mechanical anachronism!
- This has major implications for DBMS design:
  - **READ:** transfer data from disk to main memory (RAM)
  - **WRITE:** transfer data from RAM to disk
  - Both are high-cost operations, relative to in-memory operations, so must be planned carefully!

# Why not store data all in the main memory?

- ***Costs too high***

- High-end databases today in the Petabyte range
- ~ 60% of the cost of a production system is in the disks

- ***Main memory is volatile***

- We want data to be saved between runs (Obviously!)

- ***Main-memory database systems:***

- The emerging hardware enables main-memory database systems development
- Smaller size, performance optimized
- Volatility is ok for some applications

**Disks become a viable option to store data**

# Today's focus

- DBMS layers and storage hierarchy
- Disks
  - HDD
  - SSD
- Buffer management



# Disks

- Secondary storage device of choice
- Data is stored and retrieved in units called *disk blocks* or *pages*
- Unlike RAM, time to retrieve a disk page varies:
  - Depends on the type of disk (magnetic vs flash vs PCM)
  - Depends on the access pattern (sequential vs random)
- Relative placement of pages on disk has major impact on DBMS performance!

# Storage devices


- **Magnetic disks** (market value: 20B+ CHF)
  - Large capacity at low cost
  - Block level random access
  - Slow performance for random access
  - Good performance for streaming access
- **Flash memory** (market value: 32B+ CHF)
  - Capacity at intermediate cost
  - Block level random access
  - Medium performance for random writes
  - Good performance otherwise



Magnetic disks are  
60 years old!



# Magnetic disk

- Stores data magnetically on thin metallic film bonded to rotating disk of glass, ceramic, or aluminium
- **Track:** concentric circles on the disk
- **Sectors:** slice of a track; smallest addressable unit
- **Cylinder:** All the tracks under the head at a given point on all surfaces



# Magnetic disk

- Length of tracks vary across disk: outer tracks can have more sectors than inner tracks and higher bandwidth
- Disk is organized into regions (“zones”) of tracks with the same number of sectors per track
- Disk has a sector-addressable address space (sectors: 512 or 4096 bytes)
  - *Block size* is a multiple of *sector size* (fixed)
- Main operations: read/write




# Accessing a disk page

- Time to access (read/write) a disk block:
  - *Seek time* (moving arms to position disk head on track)
  - *Rotational delay* (waiting for block to rotate under head)
  - *Transfer time* (actually moving data to/from disk surface)

# Disk overhead


**Disk latency = Seek time + Rotation time + transfer time**

- **Seek:** position the head/arm over the proper track
  - To get to the track (5–15 ms)
- **Rotational delay:** Wait for desired sector to rotate under read/write head (4–8 ms)
  - Only need to wait for half a rotation on average
- **Transfer time:** Transfer a block of bits (sectors) under read/write head (25–50 usec)
- Sequential access >10x faster than random access



# Arranging pages on the disk

- “*Next*” block concept:
  - Blocks on same track, followed by
  - Blocks on same cylinder, followed by
  - Blocks on adjacent cylinder
- Blocks in a file should be arranged sequentially on disk (by “next”), to minimize seek and rotational delay.
- An important optimization: pre-fetching
- Adjacent blocks are equidistant wrt access time from starting block



**Disk block has more than one neighbor**

# Disk performance example

**Using disk efficiently:** Minimize seek time and rotational delay

Q. How do access pattern influence reads/writes to disk?

- Average seek time: 5ms
- 7200 RPM  $\Rightarrow$  Time for rotation:  $60,000 \text{ (ms/min)} / 7200 \text{ (rev/min)} \approx 8\text{ms}$
- Transfer rate of 50 MBps, with block size of 4096 bytes  
 $\Rightarrow 4096 \text{ bytes} / (50 \times 10^6) = 81.92 \times 10^{-6} \text{ sec} \approx 0.82 \text{ ms for 1 sector}$



# Disk performance example

- Read block from **random place on a disk** (random reads):
  - Seek (5ms) + Rot. delay (4ms) + Transfer (0.082ms) = 9.082 ms
  - Approx. 9ms to fetch/put data: 4096 bytes /  $9.082 \times 10^{-3} \text{ s} \approx 451 \text{ KB/s}$
- Read block from a **random place in the same cylinder on a disk**:
  - Rot. delay (4ms) + Transfer (.082ms) = 4.082 ms
  - Approx 4ms to fetch/put data: 4096 bytes /  $4.082 \times 10^{-3} \text{ s} \approx 1.03 \text{ MB/s}$
- Read next block **on the same track** (sequential reads):
  - Transfer (0.082ms): 4096 bytes /  $9.082 \times 10^{-3} \text{ s} \approx 50 \text{ MB/s}$

# Enter Flash


- Flash chips used for >20 years
- Flash evolved
  - USB keys
  - Storage in mobile devices
  - Consumer and enterprise flash disks (SSD)



- Flash in a DBMS
  - Main storage
  - Accelerator/enabler (Specialized cache, logging device)

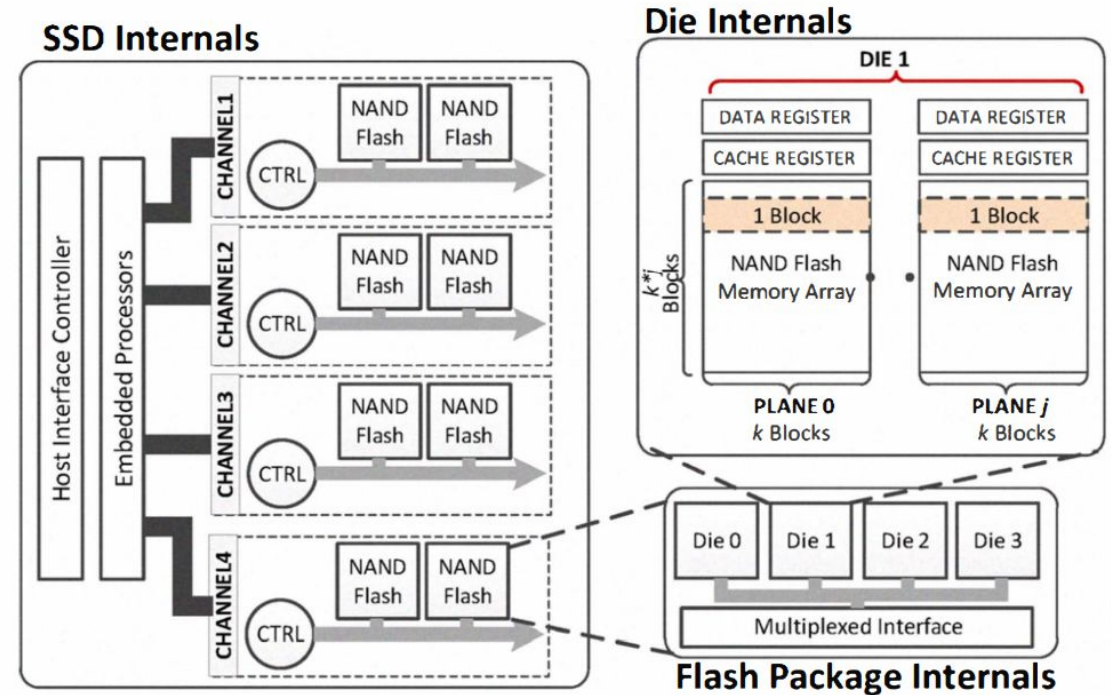
# Flash disks

- One of today's go-to solutions for fast storage
- Provides orders of magnitude higher IO performance than HDDs
- For DBMS:
  - Secondary storage *or* caching layer
  - Random reads “equally fast” as *sequential* reads
  - Data organized in *pages* (similarly to disks) and pages organized in *flash blocks*
  - *Like RAM*, time to retrieve a disk page is not related to location on flash disk



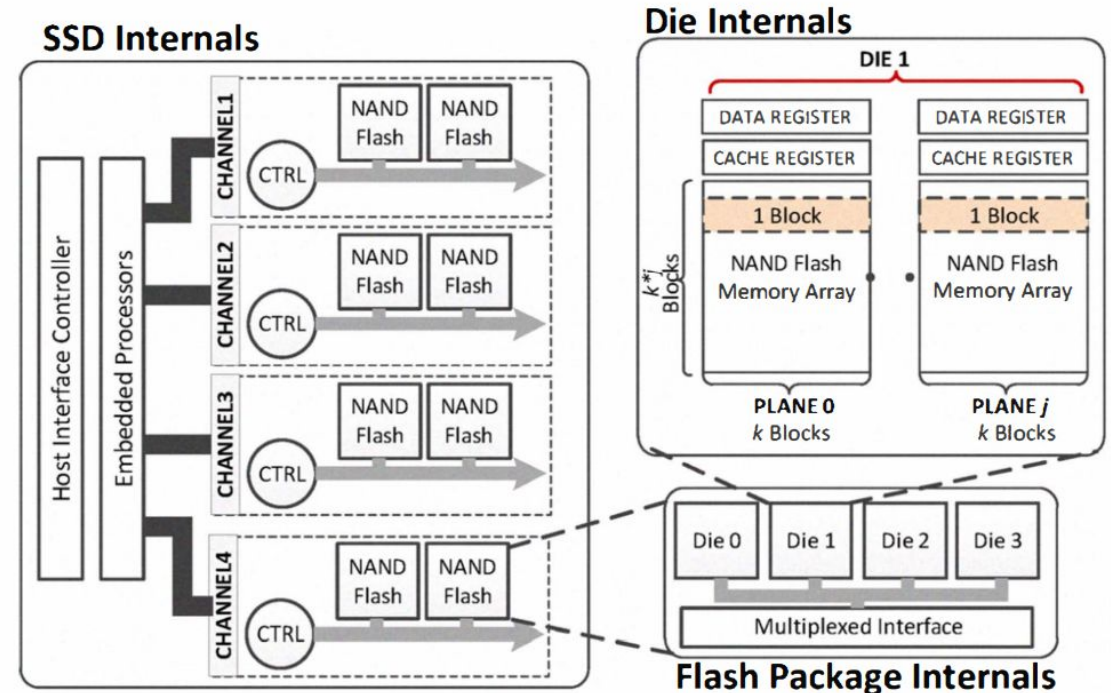
# The internals of flash disks

- Interconnected flash chips
- No mechanical limitations
- Maintain the block API – compatible with disks layout
- Internal parallelism in read/write
- Complex software driver



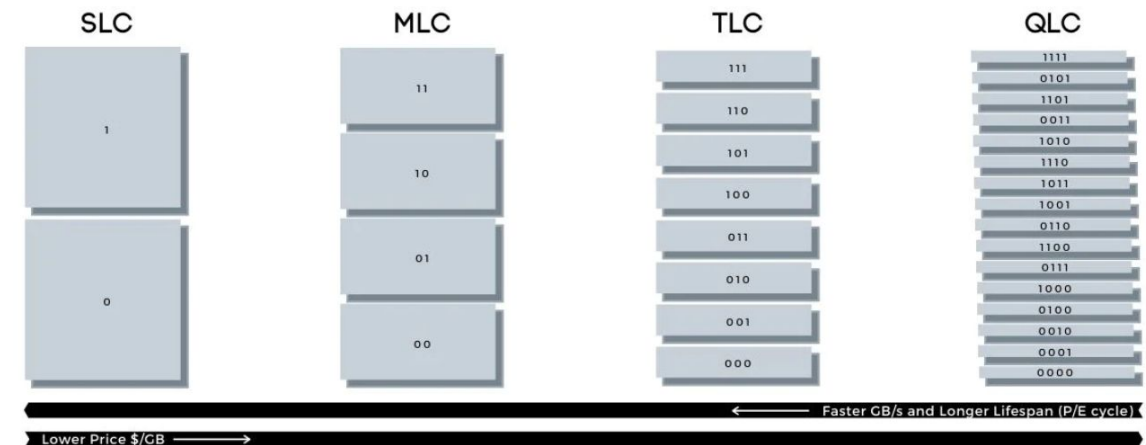
# The internals of flash disks

- SSDs have processors for IO operations
- Processors connected via channels to NAND flash package
  - Eg. 4 2GB dies (0–3)
- Each die has multiple blocks organized into planes
  - A block is 64 4 KB pages
  - Can independently access dies



# Various types of flash and supported operations

- Operations: read/write/erase
  - Read/write work at page granularity
  - Write only sets bit for cells from 1 to 0
  - Erase sets bits from 0 to 1 (block level)
- Various types of flash: SLC/MLC/TLC/QLC
  - Lifespan of devices vary from 100K–1K P/E (program/erase) cycles
  - Need to handle wear-leveling problem
    - Technique for prolonging the service life of wearable media



# Accessing a flash page

- Access time depends on
  - Device organization (internal parallelism)
  - Software efficiency (driver)
  - Bandwidth of flash packages
- Flash Translation Layer (FTL)
  - Provides similar interface as that of HDDs
  - Complex device driver (firmware)
  - Tunes performance and device lifetime


# Flash-based vs HDD-based cloud storage

## HDD

- ✓ Large – inexpensive capacity
- x Inefficient random reads

## Flash disks

- x Small – expensive capacity
- ✓ Very efficient random reads





# Rules of thumb for accessing pages ...

1. Memory access much faster than disk I/O (~ 1000x)
2. “Sequential” I/O faster than “random” I/O (~ 10x)

# Storage requirements


- **Fast:** data is there when you want it
- **Reliable:** data fetched is what you stored
- **Affordable:** won't break the bank

## ⇒ Redundant Array of Inexpensive Disks (RAID)

- Build a logical disk from (many) physical disks
- Benefits:
  - **Higher throughput** (via data “stripping”)
  - **Large MTTF** (for redundancy)
    - MTTF: Mean time to failure

# RAID 0

- File is striped across disks
- No redundancy of data (no fault tolerance)
  - Data is not mirrored on a different disk
- Provides best performance
  - **Cumulative bandwidth utilization**
- Total storage capacity: sum of capacities of all disks
- **No data security**



# Striping and reliability


Striping reduces reliability

- More disks  $\rightarrow$  higher probability of some disks failing
- $N$  disks  $\rightarrow 1/N^{\text{th}}$  mean time between failures of 1 disk

**Q. How can we improve disk reliability?**

# RAID 1

- Duplicate file blocks across storage drives
  - Deals well with disk loss
  - Does not handle corruption
- Total storage capacity: 1 disk only
- Performance:
  - Reads can be parallelized
  - Writes too, so equivalent to writing to one disk
- Expensive
- Used in critical infrastructure (sensitive information)

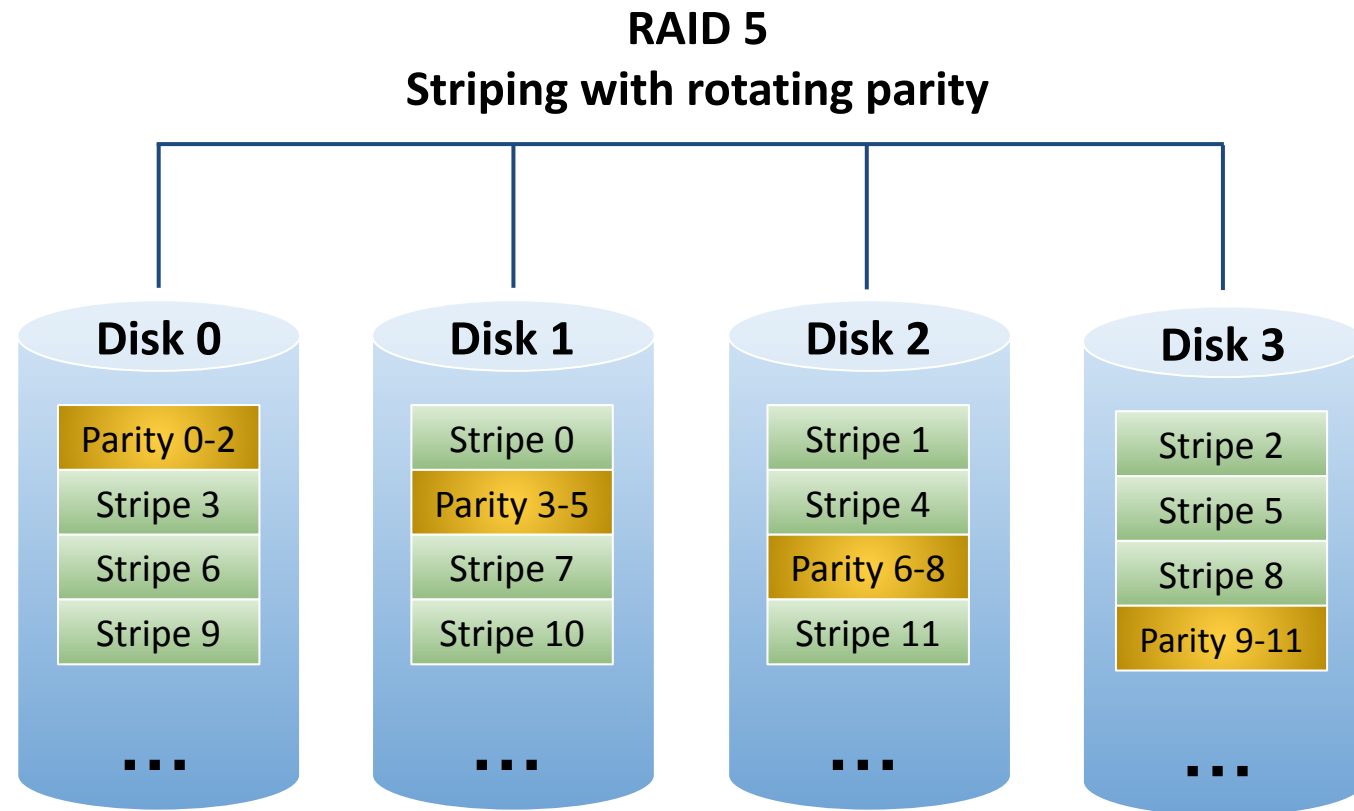


# RAID 5

- **Parity:** Another mechanism for fault tolerance.

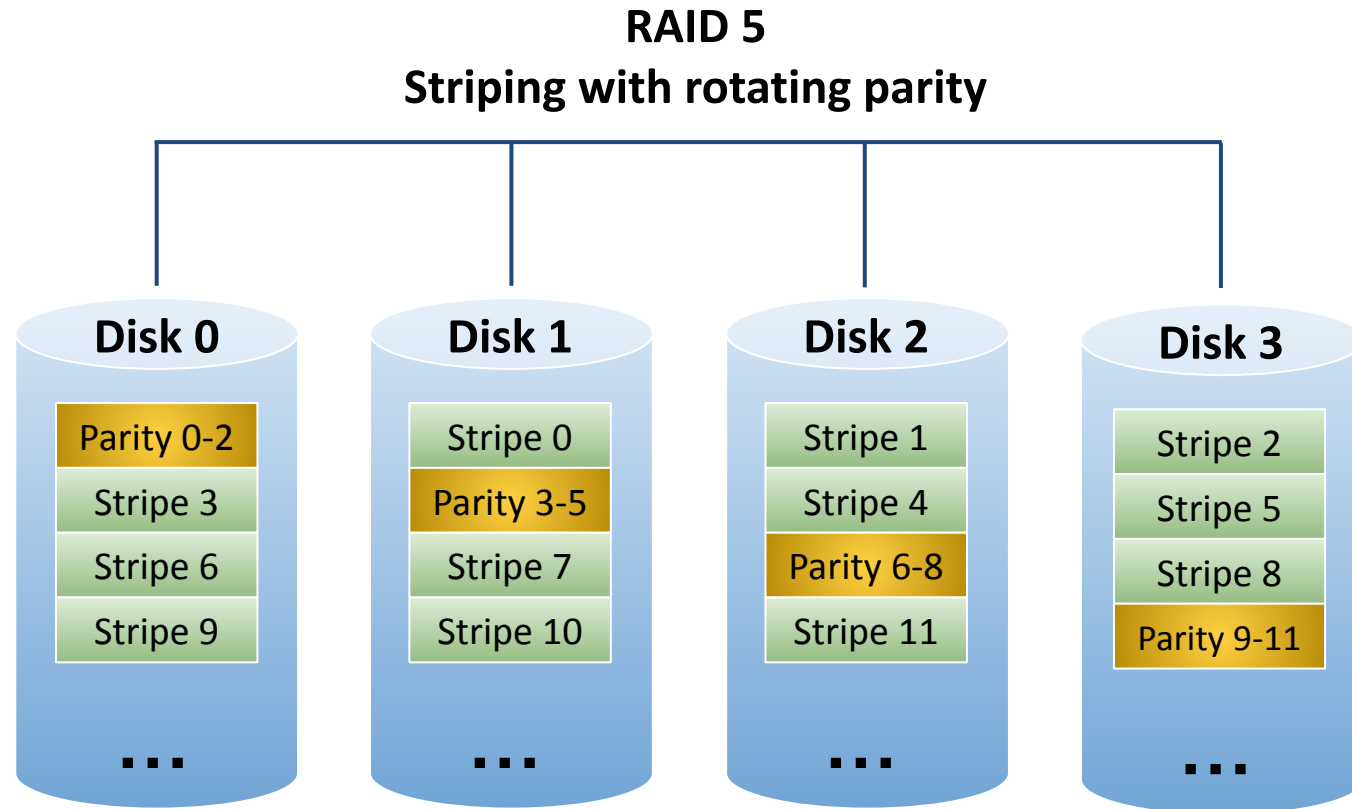
$$P_{i-j} = S_i \oplus S_{i+1} \oplus \dots \oplus S_j$$

- If one disk fails, one can reconstruct its data by XOR-ing all remaining drives



# RAID 5

- **Reliable**
  - Still works by losing one disk
- **Fast**
  - 3x seq. reads of one disk, i.e.,  $(N-1)$ ,  $N \rightarrow \# \text{ drives}$
  - Writes become complicated
- **Affordable**
- **Used:** datacenter environments



# Today's focus

- DBMS layers and storage hierarchy
- Disks
  - HDD
  - SSD
- Buffer management

Readings: Chapter 13.5

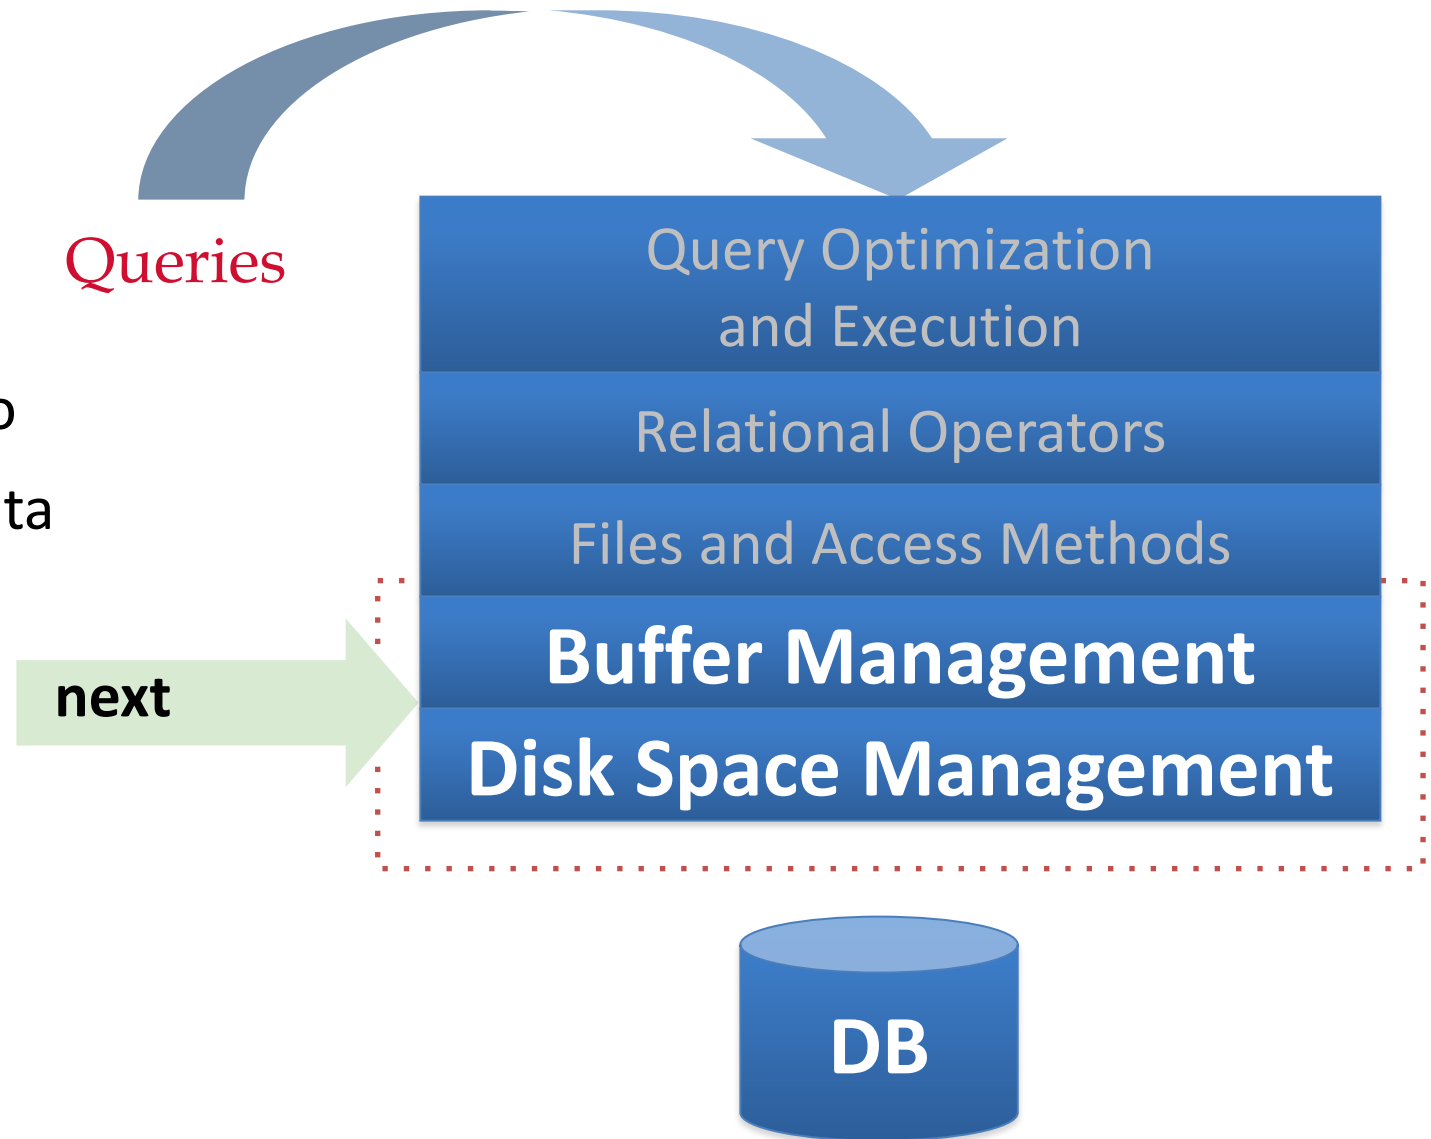


# Disk space management

- Lowest layer of DBMS software manages space on disk
- Higher levels call upon this layer to:
  - Allocate/deallocate a page
  - Read/write a page
- Best case scenario: A request for a *sequence* of pages is satisfied by pages stored sequentially on disk
  - Higher levels don't need to know if/how this is done, or how free space is managed

# Recall the big picture

DBMS relies on buffer manager to manage its memory and move data back-and-forth from disk



# Buffer management for DBMS


Q. Why not use existing OS services?

Q. How does the buffer pool **maintain** pages?

Q. What is the policy to move pages (eviction/replacement policy)?

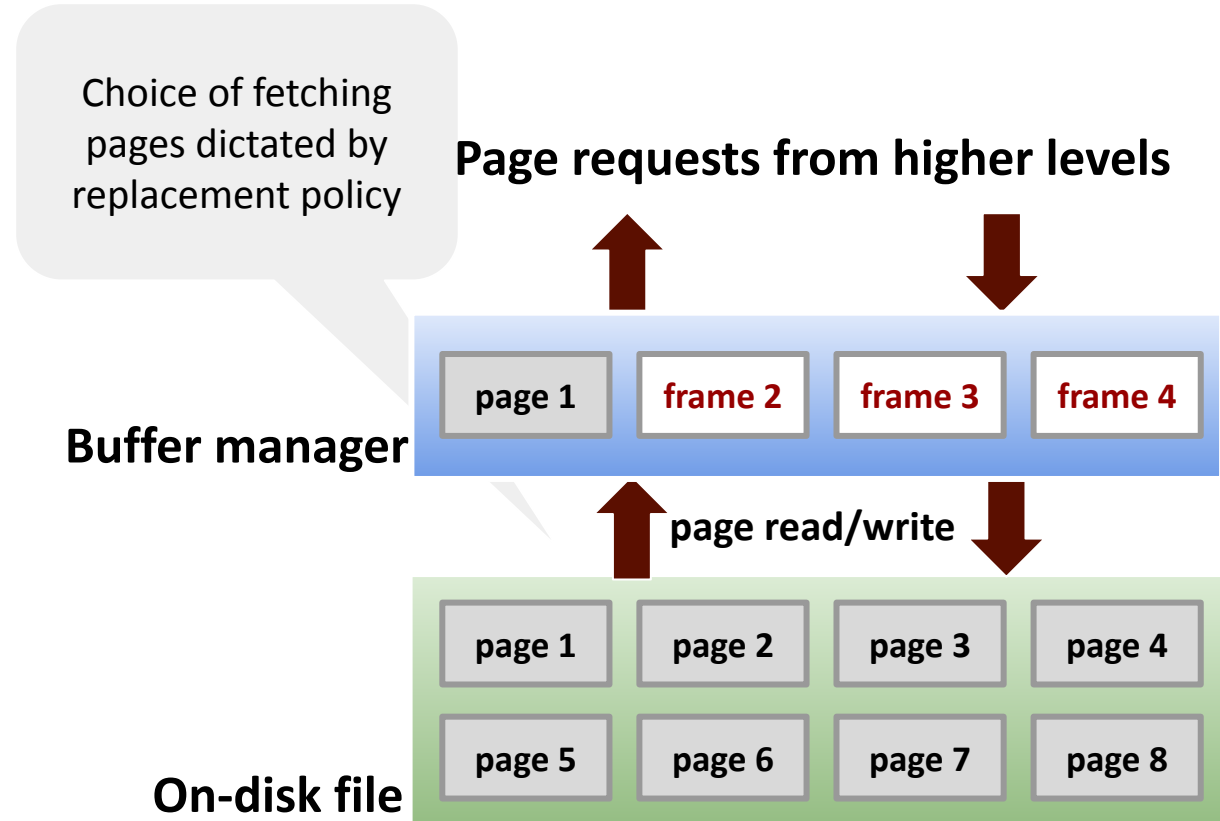
# Why not rely on OS?

- Layers of abstraction are good ... but:
  - Unfortunately, OS often **gets in the way** of DBMS
- DBMS requires to do things “its own way”
  - **Specialized prefetching**
  - **Control over buffer replacement policy**
  - **Control over flushing data to disk**
    - WAL protocol requires flushing log entries to disk
  - **Control over thread/process scheduling**
    - Convoy problem: arises when OS scheduling conflicts with DBMS locking




# Buffer management in a DBMS

- Data must be in DRAM for DBMS to operate on it
- Buffer manager hides the fact that not all data is in DRAM
  - Similar to hardware caches that hide that not all data are in the caches
- Similar to block cache, replacement policies determine which pages to keep in DRAM vs. write them back




# Buffer manager

- Reduces IO by buffering/caching
- Keeps active pages in memory
- Limited size, discards/write back pages when needed
- Also important for recovery, using logging




# Buffer manager keeps track of pages

- Memory region organized as an array of fixed-size pages
- An array entry is called a **frame**
- On requesting a page, an exact copy is placed into one of these frames
- Dirty pages are buffered in the pool to avoid disk stalls



# Buffer manager also maintains metadata for pages

- The **page table** maps page IDs to a copy of the page in buffer pool frames
  - In-memory data structure that is not stored on the disk
- Also, maintains additional metadata per-page:
  - **Dirty flag:** A page has been updated
  - **Pin/reference counter:** A marker or a counter indicating if a page is being referenced by higher levels of DBMS (within a transaction)





# Getting a page to the pool when requested ...


- If a requested page is not present in the pool, first find a frame:
  - If there is an empty frame, read the requested page into the frame
  - OR choose a frame for *replacement*  
*(only un-pinned pages are candidates)*
  - If the frame is “dirty”, write it to disk
  - Read requested page into chosen frame
- *Pin* the page and return its address
- If requests can be predicted (e.g., sequential scans)
  - Buffer manager can **pre-fetch** several pages at a time!

# More on buffer management

- Requestor of page must unpin it, and indicate whether page has been modified:
  - *dirty* bit is used for this
- A page in pool may be requested several times
  - A *pin count* is used
  - A page is a candidate for replacement iff *pin count* is 0 (*“unpinned”*)
- CC & recovery may entail additional I/O when a frame is chosen for replacement (*write-ahead log* protocol; more later)


# Buffer replacement policies

- When memory is full, some buffer pages have to be replaced
  - Discard clean pages
  - First write the dirty pages back to the disk
  - Then replaced the discarded pages with new pages
- Goals: Correctness, accuracy, speed, metadata overhead



# Buffer replacement policies

- Several policies have been explored:
  - **First in-first out (FIFO)**
  - **Least recently used (LRU)**
  - **Most frequently used (MFU)**
  - **Clock**
  - **LRU-K**
  - **2Q**



# Buffer replacement policy: FIFO

First in-first out (FIFO)

- Simple replacement policy
- Keeps those pages that were most recently added to the buffer pool
- Maintain a linked list of buffer frames
- Insert at the end, remove from the head
- “Old” pages are removed first

**Does not retain frequently-used pages (or locality)**

# Buffer replacement policy: LRU

Least recently used (LRU)

- Simple replacement policy
- Maintain a timestamp of when each page was last accessed (*unpinned*)
- When the DBMS needs to evict a page, select the one with **oldest access timestamp**
  - Keep the pages in sorted order to reduce the search time on eviction
  - Doubly-linked list stores the buffer frames
  - Remove from the head
  - When a frame is unpinned, move it to the end of the list
  - “Hot” pages are retained in the buffer

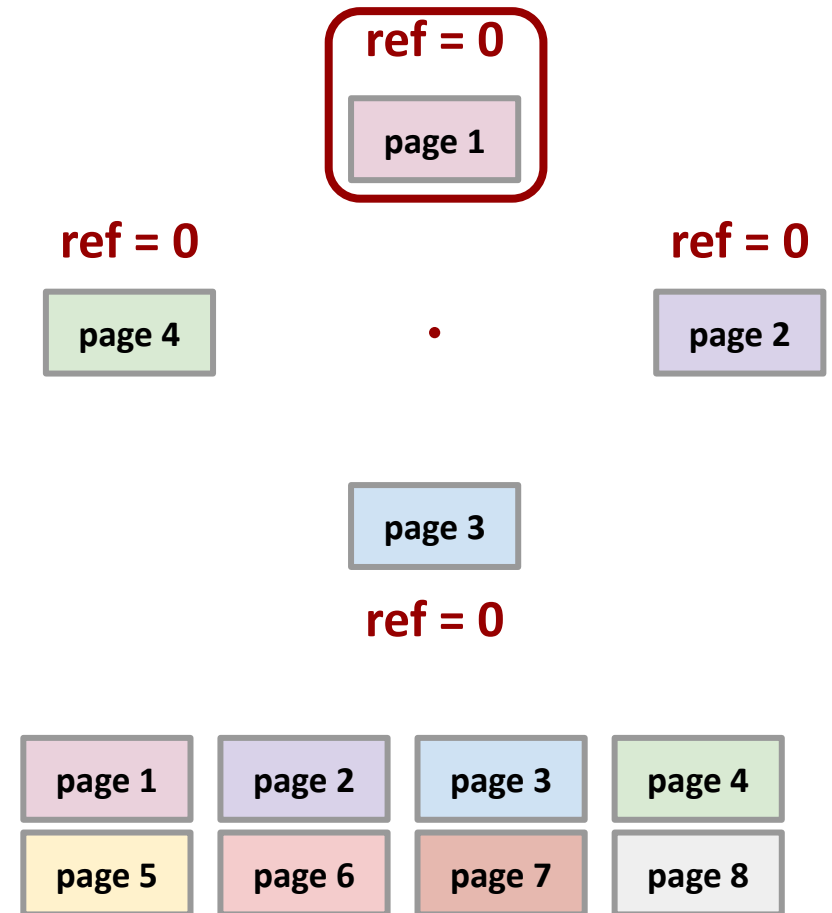
One of the most widely-used policy

# Buffer replacement policy: Clock

- An approximation of LRU that does not need a separate timestamp per page
  - Each page has a **reference bit**
  - When a page is accessed, set it to 1
- Organize the pages in a circular buffer with a “clock hand”
  - Upon sweeping, check if a page’s bit is set to 1
  - If yes, set to zero
  - If no, then evict

# Buffer replacement policy: Clock

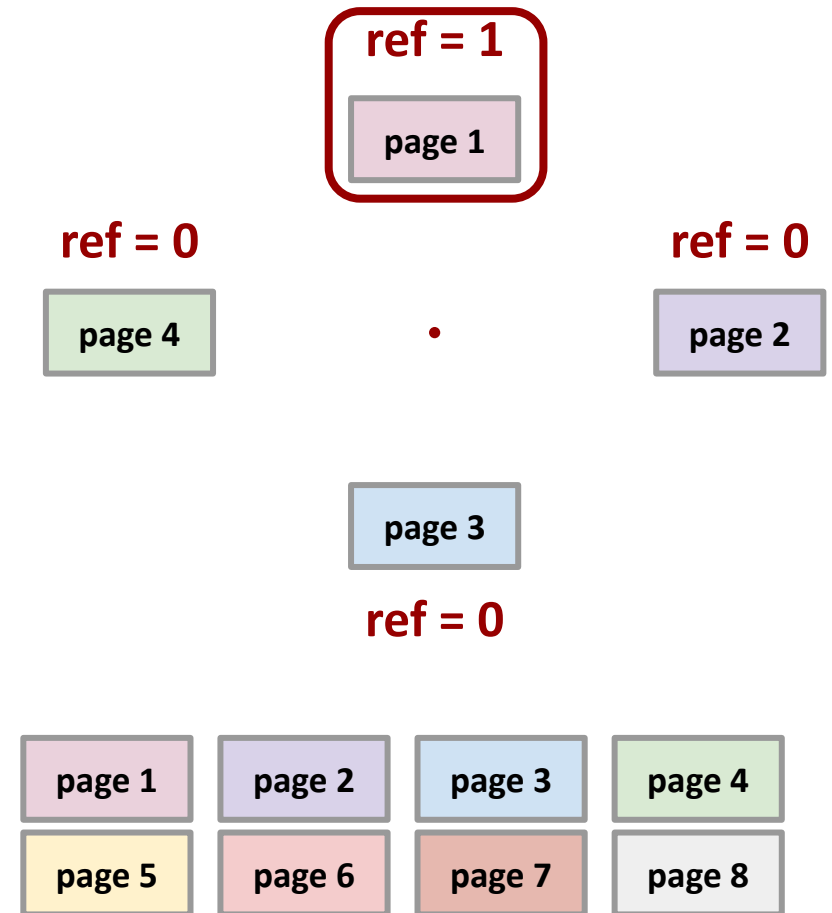
- Replacing pages:  
do {  
    if (pincount == 0 && ref bit is off)  
        choose **current** page for replacement;  
    else if (pincount == 0 && ref bit is on)  
        turn off ref bit;  
    advance current frame;  
} until a page is chosen for replacement;





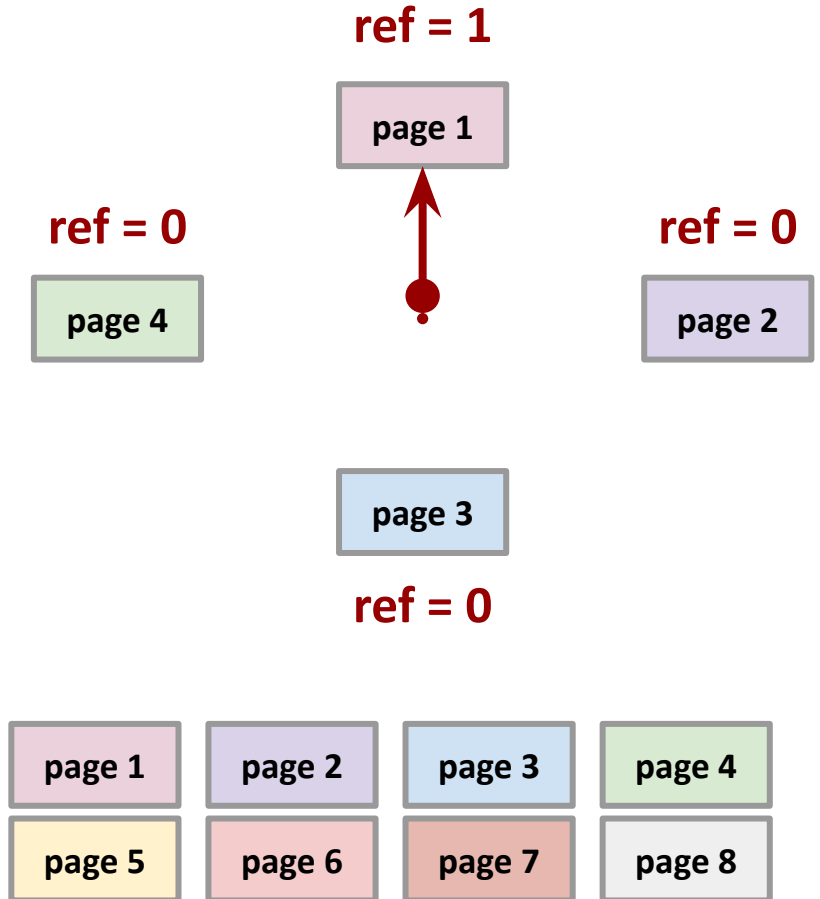
# Buffer replacement policy: Clock

- Replacing pages:  
do {  
    if (pincount == 0 && ref bit is off)  
        choose **current** page for replacement;  
    else if (pincount == 0 && ref bit is on)  
        turn off ref bit;  
    advance current frame;  
} until a page is chosen for replacement;



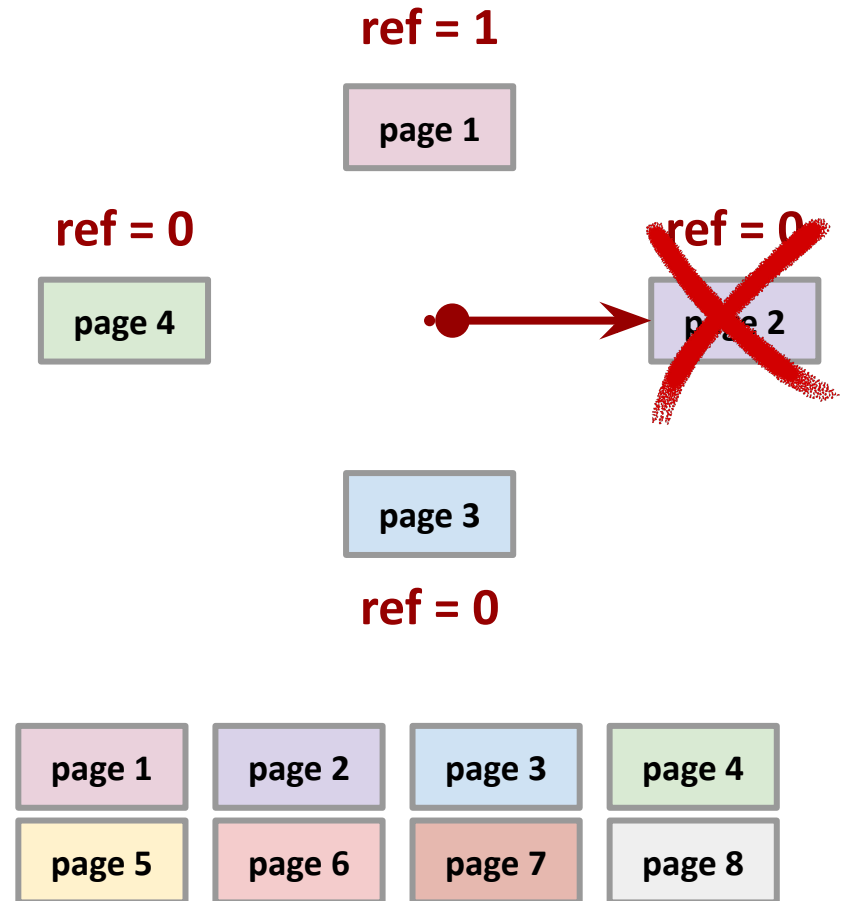
# Buffer replacement policy: Clock

- Replacing pages:  
do {  
    if (pincount == 0 && ref bit is off)  
        choose **current** page for replacement;  
    else if (pincount == 0 && ref bit is on)  
        turn off ref bit;  
    advance current frame;  
} until a page is chosen for replacement;



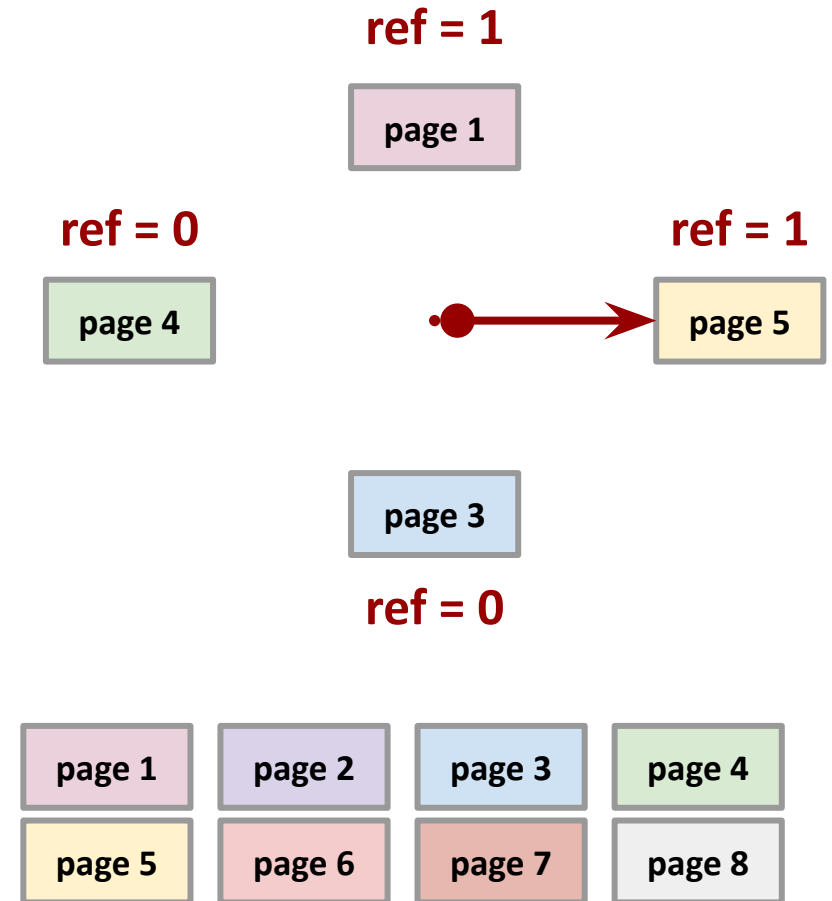
# Buffer replacement policy: Clock

- Replacing pages:  
do {  
    if (pincount == 0 && ref bit is off)  
        choose **current** page for replacement;  
    else if (pincount == 0 && ref bit is on)  
        turn off ref bit;  
    advance current frame;  
} until a page is chosen for replacement;



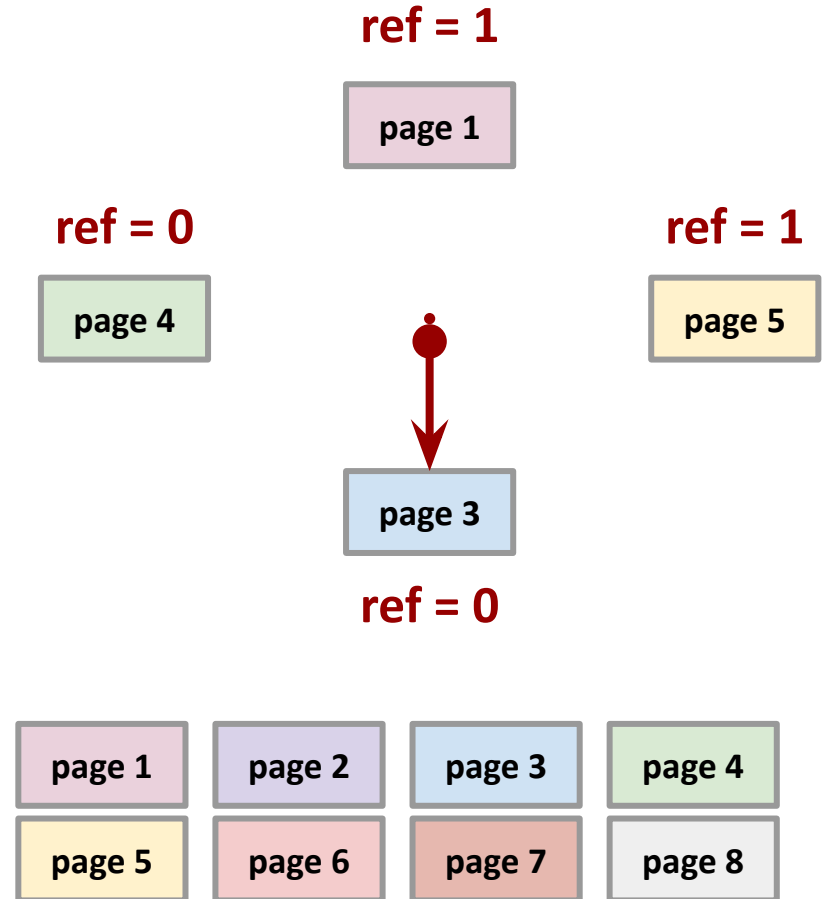
# Buffer replacement policy: Clock

- Replacing pages:  
do {  
    if (pincount == 0 && ref bit is off)  
        choose **current** page for replacement;  
    else if (pincount == 0 && ref bit is on)  
        turn off ref bit;  
    advance current frame;  
} until a page is chosen for replacement;





# Buffer replacement policy: Clock

- Replacing pages:  
do {  
    if (pincount == 0 && ref bit is off)  
        choose **current** page for replacement;  
    else if (pincount == 0 && ref bit is on)  
        turn off ref bit;  
    advance current frame;  
} until a page is chosen for replacement;



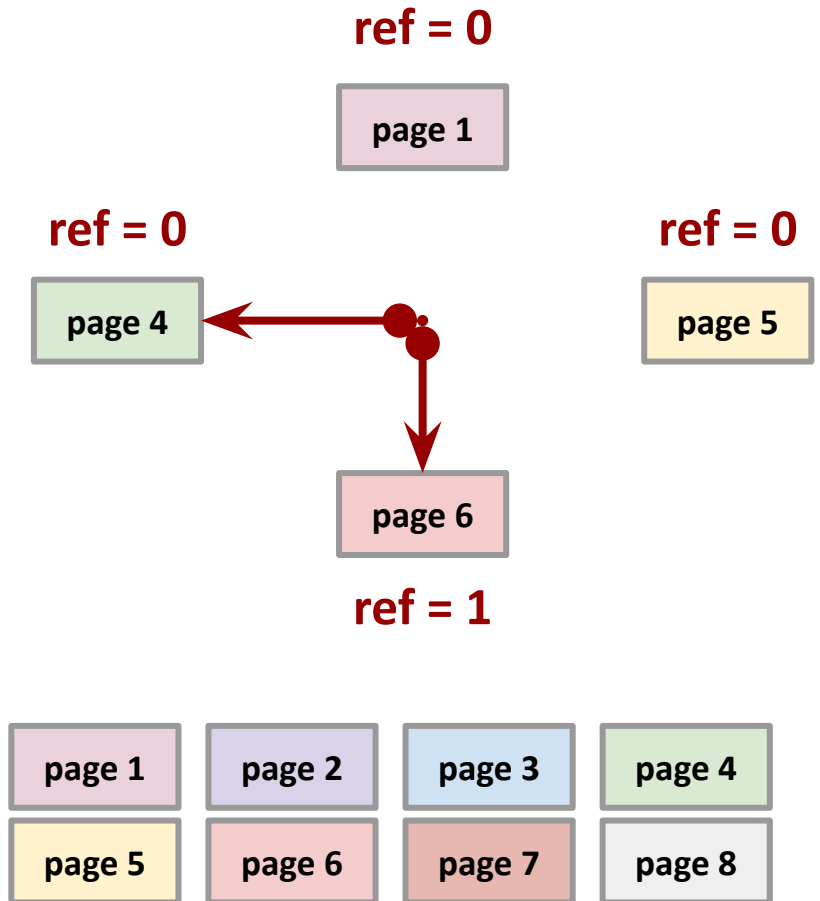
# Buffer replacement policy: Clock

- Replacing pages:  
do {  
    if (pincount == 0 && ref bit is off)  
        choose **current** page for replacement;  
    else if (pincount == 0 && ref bit is on)  
        turn off ref bit;  
    advance current frame;  
} until a page is chosen for replacement;




# Buffer replacement policy: Clock

- Replacing pages:  
do {  
    if (pincount == 0 && ref bit is off)  
        choose **current** page for replacement;  
    else if (pincount == 0 && ref bit is on)  
        turn off ref bit;  
    advance current frame;  
} until a page is chosen for replacement;



# Problem with LRU and clock

- LRU and clock are susceptible to **sequential flooding** (no scan resistance)
  - A query performs a sequential scan that reads every page
  - This **pollutes** the buffer pool with pages that are read once and then never again
  - **# buffer frames < # pages in file** means each page request causes an I/O
  - LRU-K and 2Q minimize this issue





# Buffer replacement policy: LRU-K

Least recently used - K (LRU)

- Track the history of last K references to each page as timestamps and compute the interval between subsequent accesses
- DBMS uses this history to estimate the next time that page is going to be accessed
- Becomes classic LRU when  $K = 1$
- LRU-K is scan resistant

# Buffer replacement policy: 2Q

Maintains two queues (FIFO and LRU)

- Some pages are accessed only once (eg, sequential scan)
- Some pages are hot and accessed frequently
- Maintain separate lists for those pages
- Scan resistant policy

Approach:

1. First maintain all pages in FIFO queue
2. When a page that is currently in FIFO is referenced again, upgrade it to LRU queue
3. Prefer evicting pages from FIFO queue

**Hot pages are in LRU, read-once pages are in FIFO: good strategy for DBMS**

# Summary

- Disks provide cheap, non-volatile storage
  - Random access, but cost depends on location of page on disk; important to arrange data sequentially to minimize *seek* and *rotation* delays
- Buffer manager brings pages into RAM
  - Page stays in RAM until released by requestor
  - Written to disk when frame chosen for replacement (which is sometime after requestor releases the page)
  - Choice of frame to replace based on *replacement policy*
  - Good to *pre-fetch* several pages at a time

# Buffer manager implementation

- The hash table uses latches for mutual exclusion
- PageNo: The page number
- Latch: A read/write lock to protect the page
- LSN: log sequence number

(a unique identifier associated with every record in a transactional log)

- State: Clean/dirty/new created etc.
- Data: The actual data contained on the page

Hashtable



Buffer frames

